

Database Systems: Project 2 Report

Marshall Cooper - Andrea Piccione

In this project, we implemented a two-phase CUBE operator on Spark RDDs as well as a similarity ClusterJoin. For our CUBE operator, we compared its performance in a variety of tests versus a more naïve MapReduce implementation with only a single phase. We tested over three sizes of the same data set and varied the number of reducers as well as the CUBE dimension $|D|$ (*number of columns used for grouping*). We intentionally chose CUBE columns to allow for some aggregation to occur, meaning that we did not, for example, choose any column that uniquely identifies the data set (*such as a primary key or comment field*). The following summarizes the results of testing:

Runtime (s) for CUBE algorithms on 6008 tuples, 16 reducers

	Dimension $ D $ of CUBE					
	1	2	3	4	5	6
Naïve	0.0767	0.0647	0.2106	0.3599	0.7863	2.6452
2-Phase	0.0625	0.0694	0.3314	0.4537	0.8956	2.0827

Runtime (s) for 2 Dimension CUBE, 16 reducers

	Number of tuples		
	6008	600572	6001171
Naïve	0.2849356	0.6053815	9.1548128
2-Phase	0.3876628	0.6597307	6.1207095

Runtime (s) for 3 Dimension CUBE on 600572 tuples

	Number of reducers				
	1	4	8	16	32
Naïve	8.3122	3.4269	2.2873	1.9909	1.8037
2-Phase	23.788	10.378	4.64	3.1343	2.1077

The first interpretation that we make for these results is that in simple cases, with a low $|D|$ and a small size data set, the naïve operator substantially outperforms the two-phase optimization. For the two-phase algorithm, there is an additional mapper side combine phase that is meant to lower the workload on reducers, but we believe that the extra execution of this phase, for small inputs, is not offset by the reduction in workload further down the stream. We observe that on larger data sets, the two-phase algorithm suffers less in performance than the naïve implementation. More drastically, we note that the two-phase algorithm deals substantially better with higher $|D|$. This appears to be the most drastic performance benefit for the two-phase algorithm, as even on the small data set, it handled larger CUBE dimensions more easily. Finally, we note that increasing the number of reducers seems to benefit both algorithms up to a point. For the case of the two-phase algorithm, which has multiple reduction steps, we see that having a small number of reducers is especially penalizing. Prior to testing, we had the intuition that the naïve algorithm would perform better on less computationally intensive test cases because it is a simpler algorithm with fewer steps. Ultimately we were surprised that the optimized version only outperformed the naïve algorithm when either the number of tuples or the dimension of the CUBE was very large. Nonetheless, this indicates that is a more scalable algorithm.

In the second task of this project, we implemented a similarity join operator using an optimized ClusterJoin algorithm [1]. We carried out a battery of performance tests that aim at assessing the validity and the scalability

of the ClusterJoin algorithm in comparison with a reference implementation which takes advantage of the Cartesian product to perform similarity checks among each possible combination of strings in the chosen column of the table.

We performed one series of tests to see the influence the number of anchors had on the ClusterJoin, and additionally performed a series of tests to see how the ClusterJoin would compare versus the Cartesian product algorithm on differently sized samples of the input data. We sampled 10%, 20%, 50% and 100% of the original dataset (10k inputs) for testing. For all tests, the distance threshold was set to 2 as a constant. We made the choice to control this feature because it doesn't directly affect the time elapsed for the computation of the result.

ClusterJoin Runtime (s)

		Number of Anchors					
		1	2	4	8	16	32
Input size	1000	25.2808	24.7078	14.692	21.5589	23.3682	15.0509
	2000	50.6235	60.8314	53.2724	50.2225	46.8363	51.9872

Runtime (s) Comparison between ClusterJoin and Cartesian Product-based Join

		Input size							
		1k			2k			5k	10k
		Number of Anchors			Number of Anchors			Anchors	Anchors
		1	4	16	1	4	16	4	4
ClusterJoin		25.280752	14.692032	23.368197	50.204785	53.255414	53.463841	295.05756	1136.5164
Cartesian		13.99230057			46.83632267			358.26465	1201.0249

We can first notice that, as a general trend for low dimensional input sizes, increasing the number of anchors makes the ClusterJoin algorithm faster. This was an expected behavior predicted by the authors of the algorithm [1], who state: "for low dimension spatial data, having a larger number of anchors will likely make the "closest" anchor even closer, thus improving the effectiveness of our filtering procedure". On the other hand, that authors also state that "in high dimensional space, adding more anchors does not significantly reduce the average distance to the "closest" anchor." This is reflected by our observed results that show how the performance improvement becomes significantly less evident already by sampling 20% of the original input dataset.

Varying the size of the dataset provided us with an interesting overview of the scalability capabilities of our implementation. In fact, even though the Cartesian implementation evidently outperforms the ClusterJoin algorithm for small data sizes, the efficiency of the optimized algorithm is noticeable and can be appreciated for larger input sizes where it is definitely faster than computing the Cartesian product for each string pair.

Therefore, we believe these observation support the scalability of the ClusterJoin algorithm. We also noticed that the difference in performance between the algorithms, even on small data sets, was very small if a reasonable number of anchors was used.

In conclusion, both algorithms we implemented showed clear signs of increased scalability over their more simplistic counterparts. In the case of the CUBE operator, we were surprised how large the scale of parameters had to be in order for the optimized algorithm to outperform its counterpart. In the case of the ClusterJoin implementation, we were surprised that the ClusterJoin was proportionally competitive for even small inputs and outperformed its counterpart even for a relatively small (10k) selection of strings.

Reference

[1] Sarma, A., He, Y., Chaudhury, S. ClusterJoin: A Similarity Joins Framework using Map-Reduce, 2014